

# MultipleModels4-28-mod

April 30, 2025

```
[1]: # Hybrid Linear Regression + Random Forest Regression

# Generic inputs for most ML tasks
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

import warnings # getting lots of user warnings from sklearn
warnings.filterwarnings("ignore", category=UserWarning)
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import LabelEncoder
import datetime
from sklearn import tree
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score

pd.options.display.float_format = '{:,.2f}'.format

# setup interactive notebook mode
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
from IPython.display import display, HTML

# Initialize lists to collect model performance
model_names = []
model_maes = []
model_mabs = []

[2]: data_set = pd.read_csv('MultipleSources-2019-2025-Ascending.csv')
data_set.head()
data_set.tail()
```

```
[2]:      Date      SP500  Futures    Nikkei    FTSE      DAX
0  04/03/19  2,873.40  2,879.75  21,713.21  7,418.28  11,954.40
1  04/04/19  2,879.39  2,882.75  21,724.95  7,401.94  11,988.01
2  04/05/19  2,892.74  2,896.00  21,807.50  7,446.87  12,009.75
3  04/08/19  2,895.77  2,898.25  21,761.65  7,451.89  11,963.40
4  04/09/19  2,878.20  2,882.50  21,802.59  7,425.57  11,850.57
```

```
[2]:      Date      SP500  Futures    Nikkei    FTSE      DAX
1516  04/14/25  5,405.97  5,440.75  33,982.36  8,134.34  20,954.83
1517  04/15/25  5,396.63  5,386.25  34,267.54  8,249.12  21,253.70
1518  04/16/25  5,275.70  5,305.75  33,920.40  8,275.60  21,311.02
1519  04/17/25  5,282.70  5,312.75  34,730.28  8,275.66  21,205.86
1520  04/21/25  5,158.20  5,184.00  34,279.92      NaN      NaN
```

```
[3]: data_set.columns
data_set['Date'] = pd.to_datetime(data_set['Date'])
# get day of week
data_set['day_of_week'] = data_set['Date'].dt.dayofweek
data_set['Date'] = data_set['Date'].astype(np.int64) // 10**9
data_set.tail()
```

```
[3]: Index(['Date', 'SP500', 'Futures', 'Nikkei', 'FTSE', 'DAX'], dtype='object')
```

```
[3]:      Date      SP500  Futures    Nikkei    FTSE      DAX  day_of_week
1516  1744588800  5,405.97  5,440.75  33,982.36  8,134.34  20,954.83      0
1517  1744675200  5,396.63  5,386.25  34,267.54  8,249.12  21,253.70      1
1518  1744761600  5,275.70  5,305.75  33,920.40  8,275.60  21,311.02      2
1519  1744848000  5,282.70  5,312.75  34,730.28  8,275.66  21,205.86      3
1520  1745193600  5,158.20  5,184.00  34,279.92      NaN      NaN      0
```

```
[4]: # Reshape as required by OneHotEncoder
encoder = OneHotEncoder(handle_unknown='ignore', sparse_output=False)
day_of_week_encoded = encoder.fit_transform(data_set[['day_of_week']])
```

```
[5]: # day of week as feature
encoded_cols = encoder.get_feature_names_out(['day_of_week'])
encoded_df = pd.DataFrame(day_of_week_encoded, columns=encoded_cols,
    ↪ index=data_set.index)
data_set = pd.concat([data_set.drop(columns=['day_of_week']), encoded_df],
    ↪ axis=1)
```

```
[6]: def add_lag_features(df, columns, lags):
    """ Add lag features for columns and lag days.
    columns (list): Columns for which to create lag features
    lags (list): List of lag values (e.g., [1, 2, 3]) """
    for col in columns:
        for lag in lags:
            df[f'{col}_lag_{lag}'] = df[col].shift(lag)
```

```

return df

lag_columns = ['SP500', 'DAX', 'FTSE', 'Nikkei']
lag_days = [1, 2, 3]

data_set = add_lag_features(data_set, lag_columns, lag_days)
data_set = data_set.dropna()
data_set.head()

data_set = data_set.sort_values(by='Date').reset_index(drop=True)

```

```

[6]:
      Date      SP500  Futures   Nikkei   FTSE      DAX  day_of_week_0  \
3  1554681600  2,895.77  2,898.25  21,761.65  7,451.89  11,963.40      1.00
4  1554768000  2,878.20  2,882.50  21,802.59  7,425.57  11,850.57      0.00
5  1554854400  2,888.21  2,894.50  21,687.57  7,421.91  11,905.91      0.00
6  1554940800  2,888.32  2,891.75  21,711.38  7,417.95  11,935.20      0.00
7  1555027200  2,907.41  2,912.50  21,870.56  7,437.06  11,999.93      0.00

      day_of_week_1  day_of_week_2  day_of_week_3  ...  SP500_lag_3  DAX_lag_1  \
3              0.00              0.00              0.00  ...      2,873.40  12,009.75
4              1.00              0.00              0.00  ...      2,879.39  11,963.40
5              0.00              1.00              0.00  ...      2,892.74  11,850.57
6              0.00              0.00              1.00  ...      2,895.77  11,905.91
7              0.00              0.00              0.00  ...      2,878.20  11,935.20

      DAX_lag_2  DAX_lag_3  FTSE_lag_1  FTSE_lag_2  FTSE_lag_3  Nikkei_lag_1  \
3  11,988.01  11,954.40    7,446.87    7,401.94    7,418.28    21,807.50
4  12,009.75  11,988.01    7,451.89    7,446.87    7,401.94    21,761.65
5  11,963.40  12,009.75    7,425.57    7,451.89    7,446.87    21,802.59
6  11,850.57  11,963.40    7,421.91    7,425.57    7,451.89    21,687.57
7  11,905.91  11,850.57    7,417.95    7,421.91    7,425.57    21,711.38

      Nikkei_lag_2  Nikkei_lag_3
3      21,724.95    21,713.21
4      21,807.50    21,724.95
5      21,761.65    21,807.50
6      21,802.59    21,761.65
7      21,687.57    21,802.59

[5 rows x 23 columns]

```

```

[7]: print(data_set.columns)

```

```

Index(['Date', 'SP500', 'Futures', 'Nikkei', 'FTSE', 'DAX', 'day_of_week_0',
      'day_of_week_1', 'day_of_week_2', 'day_of_week_3', 'day_of_week_4',
      'SP500_lag_1', 'SP500_lag_2', 'SP500_lag_3', 'DAX_lag_1', 'DAX_lag_2',
      'DAX_lag_3', 'FTSE_lag_1', 'FTSE_lag_2', 'FTSE_lag_3', 'Nikkei_lag_1',
      'Nikkei_lag_2', 'Nikkei_lag_3'],
      dtype='object', name='columns')

```

```
dtype='object')
```

```
[8]: # Create target columns first
data_set['SP500_next'] = data_set['SP500'].shift(-1)
data_set['Nikkei_next'] = data_set['Nikkei'].shift(-1)
data_set['FTSE_next'] = data_set['FTSE'].shift(-1)
data_set['DAX_next'] = data_set['DAX'].shift(-1)

# drop any new rows with NaNs
data_set = data_set.dropna()
```

```
[9]: # Drop the 'Date' and target columns to get features
features = data_set.drop(columns=['Date', 'SP500_next', 'Nikkei_next', 'FTSE_next', 'DAX_next'])
targets = data_set[['SP500_next', 'Nikkei_next', 'FTSE_next', 'DAX_next']]

# Apply StandardScaler to features only
scaler = StandardScaler()
scaled_features = scaler.fit_transform(features)

# Rebuild the scaled DataFrame
data_scaled = pd.DataFrame(scaled_features, columns=features.columns)
data_scaled = pd.concat([data_scaled, targets], axis=1)
```

```
[10]: print(data_scaled.columns)
```

```
Index(['SP500', 'Futures', 'Nikkei', 'FTSE', 'DAX', 'day_of_week_0',
      'day_of_week_1', 'day_of_week_2', 'day_of_week_3', 'day_of_week_4',
      'SP500_lag_1', 'SP500_lag_2', 'SP500_lag_3', 'DAX_lag_1', 'DAX_lag_2',
      'DAX_lag_3', 'FTSE_lag_1', 'FTSE_lag_2', 'FTSE_lag_3', 'Nikkei_lag_1',
      'Nikkei_lag_2', 'Nikkei_lag_3', 'SP500_next', 'Nikkei_next',
      'FTSE_next', 'DAX_next'],
      dtype='object')
```

```
[11]: # for each stock index_next

# Parameters for sliding window
window_size = 10
train_size = 8
test_size = 2

X_full = data_scaled.drop(columns=['SP500_next', 'Nikkei_next', 'FTSE_next', 'DAX_next'])
y_full = data_scaled['SP500_next']

# Containers to hold concatenated training and testing data
X_train_list, X_test_list = [], []
y_train_list, y_test_list = [], []
```

```

# Sliding window generation
for start in range(0, len(X_full) - window_size + 1, test_size): # step by
    ↪test size
    end = start + window_size
    X_window = X_full.iloc[start:end]
    y_window = y_full.iloc[start:end]

    X_train_list.append(X_window.iloc[:train_size])
    y_train_list.append(y_window.iloc[:train_size])
    X_test_list.append(X_window.iloc[train_size:])
    y_test_list.append(y_window.iloc[train_size:])

# Concatenate all windows together to get final train-test sets
X_train = pd.concat(X_train_list, ignore_index=True)
y_train = pd.concat(y_train_list, ignore_index=True)
X_test = pd.concat(X_test_list, ignore_index=True)
y_test = pd.concat(y_test_list, ignore_index=True)

```

```

[12]: ### for Nikkei
# Parameters for sliding window
window_size = 10
train_size = 8
test_size = 2

X_full_nk = data_scaled.drop(columns=['SP500_next', 'Nikkei_next', 'FTSE_next',
    ↪'DAX_next'])
y_full_nk = data_scaled['Nikkei_next']

# Containers to hold concatenated training and testing data
X_train_list_nk, X_test_list_nk = [], []
y_train_list_nk, y_test_list_nk = [], []

# Sliding window generation
for start in range(0, len(X_full_nk) - window_size + 1, test_size): # step by
    ↪test size
    end = start + window_size
    X_window_nk = X_full_nk.iloc[start:end]
    y_window_nk = y_full_nk.iloc[start:end]

    X_train_list_nk.append(X_window_nk.iloc[:train_size])
    y_train_list_nk.append(y_window_nk.iloc[:train_size])
    X_test_list_nk.append(X_window_nk.iloc[train_size:])
    y_test_list_nk.append(y_window_nk.iloc[train_size:])

# Concatenate all windows together to get final train-test sets
X_train_nk = pd.concat(X_train_list_nk, ignore_index=True)

```

```

y_train_nk = pd.concat(y_train_list_nk, ignore_index=True)
X_test_nk = pd.concat(X_test_list_nk, ignore_index=True)
y_test_nk = pd.concat(y_test_list_nk, ignore_index=True)

```

```

[13]: ### for FTSE
# Parameters for sliding window
window_size = 10
train_size = 8
test_size = 2

# Drop target from features and extract y values
X_full_fs = data_scaled.drop(columns=['SP500_next', 'Nikkei_next', 'FTSE_next',
↪ 'DAX_next'])
y_full_fs = data_scaled['FTSE_next']

# Lists to hold the training and testing splits
X_train_list_fs, X_test_list_fs = [], []
y_train_list_fs, y_test_list_fs = [], []

# Sliding window generation loop
for start in range(0, len(X_full_fs) - window_size + 1, test_size): # step by
↪ test_size
    end = start + window_size
    X_window_fs = X_full_fs.iloc[start:end]
    y_window_fs = y_full_fs.iloc[start:end]

    X_train_list_fs.append(X_window_fs.iloc[:train_size])
    y_train_list_fs.append(y_window_fs.iloc[:train_size])
    X_test_list_fs.append(X_window_fs.iloc[train_size:])
    y_test_list_fs.append(y_window_fs.iloc[train_size:])

# Concatenate all windows together
X_train_fs = pd.concat(X_train_list_fs, ignore_index=True)
y_train_fs = pd.concat(y_train_list_fs, ignore_index=True)
X_test_fs = pd.concat(X_test_list_fs, ignore_index=True)
y_test_fs = pd.concat(y_test_list_fs, ignore_index=True)

```

```

[14]: ### For DAX
# Parameters for sliding window
window_size = 10
train_size = 8
test_size = 2

# Drop target from features and extract y values
X_full_dx = data_scaled.drop(columns=['SP500_next', 'Nikkei_next', 'FTSE_next',
↪ 'DAX_next'])
y_full_dx = data_scaled['DAX_next']

```

```

# Lists to hold the training and testing splits
X_train_list_dx, X_test_list_dx = [], []
y_train_list_dx, y_test_list_dx = [], []

# Sliding window generation loop
for start in range(0, len(X_full_dx) - window_size + 1, test_size): # step by
    test_size
    end = start + window_size
    X_window_dx = X_full_dx.iloc[start:end]
    y_window_dx = y_full_dx.iloc[start:end]

    X_train_list_dx.append(X_window_dx.iloc[:train_size])
    y_train_list_dx.append(y_window_dx.iloc[:train_size])
    X_test_list_dx.append(X_window_dx.iloc[train_size:])
    y_test_list_dx.append(y_window_dx.iloc[train_size:])

# Concatenate all windows together
X_train_dx = pd.concat(X_train_list_dx, ignore_index=True)
y_train_dx = pd.concat(y_train_list_dx, ignore_index=True)
X_test_dx = pd.concat(X_test_list_dx, ignore_index=True)
y_test_dx = pd.concat(y_test_list_dx, ignore_index=True)

```

## 1 Hybrid for S&P500

```
[15]: # create for each index
```

```

model = LinearRegression(fit_intercept = True)
model.fit(X_train, y_train)

# The following gives the R-square score
model.score(X_train, y_train)

model.feature_names_in_

```

```
[15]: LinearRegression()
```

```
[15]: 0.9959135976563152
```

```
[15]: array(['SP500', 'Futures', 'Nikkei', 'FTSE', 'DAX', 'day_of_week_0',
        'day_of_week_1', 'day_of_week_2', 'day_of_week_3', 'day_of_week_4',
        'SP500_lag_1', 'SP500_lag_2', 'SP500_lag_3', 'DAX_lag_1',
        'DAX_lag_2', 'DAX_lag_3', 'FTSE_lag_1', 'FTSE_lag_2', 'FTSE_lag_3',
        'Nikkei_lag_1', 'Nikkei_lag_2', 'Nikkei_lag_3'], dtype=object)
```

```
[16]: training_residuals = y_train - model.predict(X_train)
```

```
[17]: rf = RandomForestRegressor(random_state=50, min_samples_leaf = 3, max_features=
    ↪ "sqrt")

rf = rf.fit(X_train, training_residuals)
```

```
[18]: pred_residuals = rf.predict(X_test)
y_pred = pred_residuals + model.predict(X_test)
```

```
[19]: test_output = pd.DataFrame(y_pred, index = X_test.index, columns =
    ↪ ['pred_SP500_next'])
test_output = test_output.merge(y_test, left_index = True, right_index = True)
test_output.head()
mean_absolute_error = abs(test_output['pred_SP500_next'] -
    ↪ test_output['SP500_next']).mean()
print('Mean absolute error is ')
print(mean_absolute_error)
mean_absolute_percentage_error = abs(test_output['pred_SP500_next'] -
    ↪ test_output['SP500_next']).mean() / test_output['SP500_next'].mean()
print('Mean absolute percentage error is ')
print(mean_absolute_percentage_error)

model_names.append("Hybrid Model")
model_maes.append(mean_absolute_error)
model_mabs.append(mean_absolute_percentage_error)
```

```
[19]:
```

	pred_SP500_next	SP500_next
0	2,933.95	2,939.88
1	2,883.25	2,881.40
2	2,827.80	2,811.87
3	2,833.72	2,834.41
4	2,849.66	2,850.96

Mean absolute error is

5.575971488983003

Mean absolute percentage error is

0.0013263878965601493

```
[20]: # Training R²
train_hybrid_pred = model.predict(X_train) + rf.predict(X_train)
hybrid_r2_train = r2_score(y_train, train_hybrid_pred)
print(f"Hybrid Model R² score on training set: {hybrid_r2_train:.4f}")

# Test R² (y_pred is already calculated above)
hybrid_r2_test = r2_score(y_test, y_pred)
print(f"Hybrid Model R² score on test set: {hybrid_r2_test:.4f}")
```

Hybrid Model R² score on training set: 0.9999

Hybrid Model R² score on test set: 0.9998



## 2 Hybrid Model for Nikkei

```
[21]: model_nk = LinearRegression(fit_intercept = True)
      model_nk.fit(X_train_nk, y_train_nk)
```

```
# The following gives the R-square score
model_nk.score(X_train_nk, y_train_nk)
```

```
[21]: LinearRegression()
```

```
[21]: 0.994870239991317
```

```
[22]: training_residuals_nk = y_train_nk - model_nk.predict(X_train_nk)
```

```
[23]: rf_nk = RandomForestRegressor(random_state=50, min_samples_leaf = 3,
      ↪max_features = "sqrt")

      rf_nk = rf_nk.fit(X_train_nk, training_residuals_nk)
```

```
[24]: pred_residuals_nk = rf_nk.predict(X_test_nk)
      y_pred_nk = pred_residuals_nk + model_nk.predict(X_test_nk)
```

```
[25]: test_output = pd.DataFrame(y_pred_nk, index = X_test_nk.index, columns =
      ↪['pred_Nikkei_next'])
      test_output = test_output.merge(y_test_nk, left_index = True, right_index =
      ↪True)
      test_output.head()
      mean_absolute_error = abs(test_output['pred_Nikkei_next'] -
      ↪test_output['Nikkei_next']).mean()
      print('Mean absolute error is ')
      print(mean_absolute_error)
      mean_absolute_percentage_error = abs(test_output['pred_Nikkei_next'] -
      ↪test_output['Nikkei_next']).mean() / test_output['Nikkei_next'].mean()
      print('Mean absolute percentage error is ')
      print(mean_absolute_percentage_error)

      model_names.append("Hybrid Model")
      model_maes.append(mean_absolute_error)
      model_mabs.append(mean_absolute_percentage_error)
```

```
[25]:
```

	pred_Nikkei_next	Nikkei_next
0	22,226.04	22,258.73
1	21,557.70	21,344.92
2	21,268.97	21,191.28
3	21,048.67	21,067.23
4	21,178.56	21,188.56

Mean absolute error is

40.28634935525182  
Mean absolute percentage error is  
0.0013802821738920616

```
[26]: from sklearn.metrics import r2_score

# Predict hybrid values on training set
train_hybrid_pred_nk = model_nk.predict(X_train_nk) + rf_nk.predict(X_train_nk)
hybrid_r2_train_nk = r2_score(y_train_nk, train_hybrid_pred_nk)
print(f"Hybrid Model R2 score on training set (Nikkei): {hybrid_r2_train_nk:.4f}")

# Predict hybrid values on test set (y_pred_nk is already calculated)
hybrid_r2_test_nk = r2_score(y_test_nk, y_pred_nk)
print(f"Hybrid Model R2 score on test set (Nikkei): {hybrid_r2_test_nk:.4f}")
```

Hybrid Model R<sup>2</sup> score on training set (Nikkei): 0.9998  
Hybrid Model R<sup>2</sup> score on test set (Nikkei): 0.9998

### 3 Hybrid for FTSE

```
[27]: model_fs = LinearRegression(fit_intercept = True)
model_fs.fit(X_train_fs, y_train_fs)

# The following gives the R-square score
model_fs.score(X_train_fs, y_train_fs)
```

[27]: LinearRegression()

[27]: 0.9850209404286906

```
[28]: training_residuals_fs = y_train_fs - model_fs.predict(X_train_fs)
```

```
[29]: rf_fs = RandomForestRegressor(random_state=50, min_samples_leaf = 3,
    ↪max_features = "sqrt")

rf_fs = rf_fs.fit(X_train_fs, training_residuals_fs)
```

```
[30]: pred_residuals_fs = rf_fs.predict(X_test_fs)
y_pred_fs = pred_residuals_fs + model_fs.predict(X_test_fs)
```

```
[31]: test_output = pd.DataFrame(y_pred_fs, index = X_test_fs.index, columns =
    ↪['pred_FTSE_next'])
test_output = test_output.merge(y_test_fs, left_index = True, right_index =
    ↪True)
test_output.head()
```

```

mean_absolute_error = abs(test_output['pred_FTSE_next'] -
    ↪test_output['FTSE_next']).mean()
print('Mean absolute error is ')
print(mean_absolute_error)
mean_absolute_percentage_error = abs(test_output['pred_FTSE_next'] -
    ↪test_output['FTSE_next']).mean() / test_output['FTSE_next'].mean()
print('Mean absolute percentage error is ')
print(mean_absolute_percentage_error)

model_names.append("Gradient Boosting Regressor")
model_maes.append(mean_absolute_error)
model_mabs.append(mean_absolute_percentage_error)

```

```

[31]:
    pred_FTSE_next  FTSE_next
0          7,419.69    7,428.19
1          7,234.35    7,203.29
2          7,173.25    7,163.68
3          7,233.77    7,241.60
4          7,291.36    7,296.95

```

Mean absolute error is  
7.935565454755899  
Mean absolute percentage error is  
0.00108191623238732

```

[32]: from sklearn.metrics import r2_score

# Predict hybrid values on training set
train_hybrid_pred_fs = model_fs.predict(X_train_fs) + rf_fs.predict(X_train_fs)
hybrid_r2_train_fs = r2_score(y_train_fs, train_hybrid_pred_fs)
print(f"Hybrid Model R² score on training set (FTSE): {hybrid_r2_train_fs:.4f}")

# Predict hybrid values on test set
hybrid_r2_test_fs = r2_score(y_test_fs, y_pred_fs)
print(f"Hybrid Model R² score on test set (FTSE): {hybrid_r2_test_fs:.4f}")

```

Hybrid Model R² score on training set (FTSE): 0.9994  
Hybrid Model R² score on test set (FTSE): 0.9994

## 4 Hybrid for DAX

```

[33]: model_dx = LinearRegression(fit_intercept = True)
model_dx.fit(X_train_dx, y_train_dx)

# The following gives the R-square score
model_dx.score(X_train_dx, y_train_dx)

```

```
[33]: LinearRegression()
```

```
[33]: 0.9937239790484894
```

```
[34]: training_residuals_dx = y_train_dx - model_dx.predict(X_train_dx)
```

```
[35]: rf_dx = RandomForestRegressor(random_state=50, min_samples_leaf = 3,
    ↪max_features = "sqrt")

rf_dx = rf_dx.fit(X_train_dx, training_residuals_dx)
```

```
[36]: pred_residuals_dx = rf_dx.predict(X_test_dx)
y_pred_dx = pred_residuals_dx + model_dx.predict(X_test_dx)
```

```
[37]: test_output = pd.DataFrame(y_pred_dx, index = X_test_dx.index, columns =
    ↪['pred_DAX_next'])
test_output = test_output.merge(y_test_dx, left_index = True, right_index =
    ↪True)
test_output.head()
mean_absolute_error = abs(test_output['pred_DAX_next'] -
    ↪test_output['DAX_next']).mean()
print('Mean absolute error is ')
print(mean_absolute_error)
mean_absolute_percentage_error = abs(test_output['pred_DAX_next'] -
    ↪test_output['DAX_next']).mean() / test_output['DAX_next'].mean()
print('Mean absolute percentage error is ')
print(mean_absolute_percentage_error)

model_names.append("Hybrid Model DAX")
model_maes.append(mean_absolute_error)
model_mabs.append(mean_absolute_percentage_error)
```

```
[37]:   pred_DAX_next  DAX_next
0      12,297.62  12,315.18
1      12,091.76  12,059.83
2      11,928.44  11,876.65
3      11,986.64  11,991.62
4      12,093.70  12,099.57
```

Mean absolute error is

19.979527206330523

Mean absolute percentage error is

0.0013167639776862122

```
[38]: from sklearn.metrics import r2_score
```

```
# Predict hybrid values on training set
train_hybrid_pred_dx = model_dx.predict(X_train_dx) + rf_dx.predict(X_train_dx)
```

```

hybrid_r2_train_dx = r2_score(y_train_dx, train_hybrid_pred_dx)
print(f"Hybrid Model R2 score on training set (DAX): {hybrid_r2_train_dx:.4f}")

# Predict hybrid values on test set
hybrid_r2_test_dx = r2_score(y_test_dx, y_pred_dx)
print(f"Hybrid Model R2 score on test set (DAX): {hybrid_r2_test_dx:.4f}")

```

Hybrid Model R<sup>2</sup> score on training set (DAX): 0.9998

Hybrid Model R<sup>2</sup> score on test set (DAX): 0.9998

```

[39]: def predict_index(linear_model, rf_model, target_col):
    input_dict = {}
    # Use only feature columns (excluding all targets)
    for col in features.columns:
        input_dict[col] = base_row[col]

    # Update day of week features
    current_date = pd.to_datetime(datetime.today().date()) + pd.
    ↪Timedelta(days=1)
    dow = current_date.weekday()
    for i in range(5):
        input_dict[f'day_of_week_{i}'] = 1 if i == dow else 0

    input_df = pd.DataFrame([input_dict])
    scaled_input = scaler.transform(input_df)

    # Hybrid prediction
    linear_pred = linear_model.predict(scaled_input)
    rf_residual = rf_model.predict(scaled_input)
    return linear_pred[0] + rf_residual[0]

```

```

[40]: from datetime import datetime, timedelta
    # Base setup
    base_row = data_set.iloc[-1].copy()
    today = datetime.today().date()

    # Storage for results
    forecast_results = []

    #model.feature_names_in_

    required_features = model.feature_names_in_

    # Loop through today + next 3 days
    for step in range(4):
        current_date = pd.to_datetime(today + timedelta(days=step))
        #current_date = pd.to_datetime(datetime.date.today()) + pd.Timedelta(days=1)

```

```

dow = current_date.weekday()
dow_features = {f'day_of_week_{i}': int(i == dow) for i in range(5)}

# Build input_dict
input_dict = {}
for col in required_features:
    input_dict[col] = dow_features[col] if col in dow_features else
↪base_row[col]

# Scale and predict
input_df = pd.DataFrame([input_dict])[required_features]
#scaled_input_df = pd.DataFrame(scaler.transform(input_df),
↪columns=input_df.columns)

pred_sp500 = predict_index(model, rf, "SP500")
pred_nikkei = predict_index(model_nk, rf_nk, "Nikkei")
pred_ftse = predict_index(model_fs, rf_fs, "FTSE")
pred_dax = predict_index(model_dx, rf_dx, "DAX")

# Store results
forecast_results.append({
    "Date": current_date.date(),
    "S&P500": round(pred_sp500, 2),
    "Nikkei": round(pred_nikkei, 2),
    "FTSE": round(pred_ftse, 2),
    "DAX": round(pred_dax, 2),
})

# Update the base_row
base_row["SP500"] = pred_sp500
base_row["Nikkei"] = pred_nikkei
base_row["FTSE"] = pred_ftse
base_row["DAX"] = pred_dax

# get our predictions
for pred in forecast_results:
    print(f" Predictions for {pred['Date']}")
    print(f"S&P500 : {pred['S&P500']}")
    print(f"Nikkei : {pred['Nikkei']}")
    print(f"FTSE : {pred['FTSE']}")
    print(f"DAX : {pred['DAX']}")

```

Predictions for 2025-04-30  
 S&P500 : 5387.24  
 Nikkei : 33656.39  
 FTSE : 8159.58  
 DAX : 21203.16

```

    Predictions for 2025-05-01
S&P500 : 5455.5
Nikkei : 33706.35
FTSE : 8085.93
DAX : 21103.83
    Predictions for 2025-05-02
S&P500 : 5487.6
Nikkei : 33860.45
FTSE : 8011.42
DAX : 21015.28
    Predictions for 2025-05-03
S&P500 : 5503.17
Nikkei : 34069.75
FTSE : 7935.51
DAX : 20937.35

```

```

[41]: dates = []
      prices = []
      for pred in forecast_results:
          dates.append(pred['Date'])
          prices.append(pred['S&P500'])
      # Plot
      import matplotlib.pyplot as plt
      plt.figure(figsize=(10, 5))
      plt.plot(dates, prices, marker='o', linestyle='-', color='green', label='S&P500 Closing Price Predictions')
      plt.title('S&P500')
      plt.xlabel('Date')
      plt.ylabel('Predicted Price')
      plt.grid(True)
      plt.legend()
      plt.xticks(rotation=45)
      plt.tight_layout()
      plt.show()

```

[41]: <Figure size 1000x500 with 0 Axes>

[41]: [<matplotlib.lines.Line2D at 0x7f56ab2ef390>]

[41]: Text(0.5, 1.0, 'S&P500')

[41]: Text(0.5, 0, 'Date')

[41]: Text(0, 0.5, 'Predicted Price')

[41]: <matplotlib.legend.Legend at 0x7f56ab05e510>

```
[41]: (array([20208. , 20208.5, 20209. , 20209.5, 20210. , 20210.5, 20211. ]),
      [Text(20208.0, 0, '04-30 00'),
       Text(20208.5, 0, '04-30 12'),
       Text(20209.0, 0, '05-01 00'),
       Text(20209.5, 0, '05-01 12'),
       Text(20210.0, 0, '05-02 00'),
       Text(20210.5, 0, '05-02 12'),
       Text(20211.0, 0, '05-03 00')])
```

